

Workflow Overview

1. Ingestion (Bronze Layer)

- Pull Fingrid hydro power API data (JSON/CSV) using Spark notebook.
- Store directly in Fabric Lakehouse (bronze tables).

2. Cleaning & Transformation (Silver Layer)

- Normalize timestamps, units, missing values.
- Join with **hydro plant metadata** (capacity, type).

3. Analytics (Gold Layer)

- Aggregations: hourly/daily averages, regional summaries.
- Calculate load factor = actual output / capacity.
- Store as clean fact tables for dashboards.

4. Streaming Simulation

- Use Fabric Eventstream (or Python generator) to push real-time power readings.
- Spark Structured Streaming job consumes → stores into Lakehouse.

5. CI/CD

- GitHub repo with all notebooks & SQL scripts.
- GitHub Actions workflow runs unit tests (Spark job validation, SQL checks).
- Deploy updated pipelines automatically into Fabric workspace.

6. Visualization (Power BI)

- Simple dashboard with:
 - Hydro production trends (hourly, daily)
 - Regional breakdowns
 - Anomaly detection (e.g., sudden dips in output)

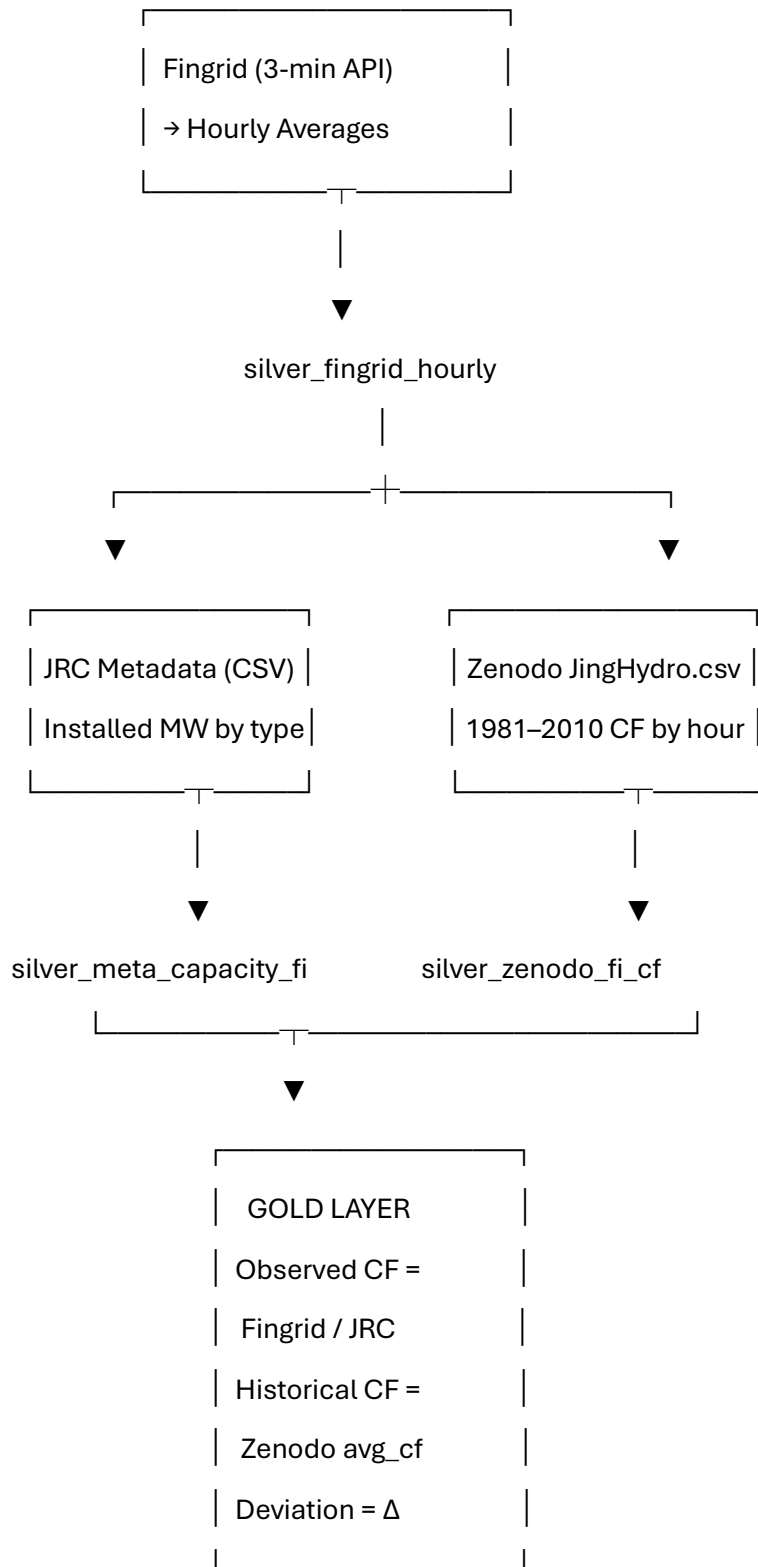
Dataset	What It Tells	Time Dimension	Granularity	Purpose
Fingrid API	Actual real-time generation (MW)	Minutes	National total	Observed performance
Zenodo	Modeled long-term capacity factors	Hourly (1981–2010)	Country/bidding zone	Historical baseline
Metadata	Installed plant capacity, type, location	Static	Plant-level	Reference and scaling factor

Layer	Dataset	Granularity	What It Brings	How It Connects
Silver	silver_fingrid_hourly	Hourly	Real generation (MW)	Used as numerator in Observed CF
Silver	silver_meta_capacity_fi	Static	Installed capacity (MW) by type	Denominator for Observed CF
Silver	silver_zenodo_fi_cf	Hourly (1981–2010)	Historical capacity factors by hour/month	Baseline to compare observed CF
Gold	—	Hourly (aligned)	Joins all three → Observed CF, Historical CF, and Deviation	Final KPI output

So:

$$\text{Observed CF} = \frac{\text{Fingrid Hourly MW}}{\text{Installed Capacity (MW)}}$$

$$\text{Deviation} = \text{Observed CF (2025)} - \text{Historical CF (1981–2010 average)}$$



Hydropower Performance Analysis — From Data Pipeline to Power BI Dashboard

This project demonstrates the end-to-end development of a modern data-engineering pipeline and analytics model for assessing Finland's hydropower efficiency against long-term climatic baselines. The objective was to integrate real-time operational data with historical hydrological benchmarks, transforming raw feeds into actionable insights through Microsoft Fabric and Power BI.

The process began with data ingestion at the **Bronze layer**, where three complementary sources were collected.

1. **Fingrid API** provided real-time electricity generation from Finnish hydro power plants at 3-minute intervals. These records reflect the country's live production output in megawatts (MW).
2. **Zenodo's JingHydro dataset** offered nationally aggregated hourly capacity-factor data for 30 climatic years (1981-2010), serving as a long-term baseline of hydro potential under historical conditions.
3. **The JRC Hydro-Power-Plant Database** supplied metadata on installed capacities and plant typologies (run-of-river, storage, and pumped-storage), giving crucial context for calculating efficiency metrics.

At the **Silver layer**, each dataset was cleaned and standardized. Fingrid data were resampled from 3-minute intervals to hourly averages and stored as `silver_fingrid_hourly`. Zenodo's European-wide data were filtered for Finland and separated into run-of-river and storage components (`silver_zenodo_fi_cf`). Metadata were aggregated to produce Finland's total installed hydro capacity (`silver_meta_capacity_fi`).

The **Gold layer** combined these curated sources into analytics-ready tables. Observed capacity factors were computed by dividing actual generation (Fingrid) by total installed capacity (metadata). Historical capacity factors were drawn from Zenodo's long-term averages. A deviation metric captured how current operational performance diverged from climatological expectations. Monthly aggregates produced two key tables:

- `gold_cf_deviation_monthly` — observed CF, historical CF, and deviation.
- `gold_zenodo_baseline_monthly` — 30-year mean and standard deviation of historical CFs.

Finally, results were presented through **Power BI visualizations** connected directly to the Fabric Lakehouse in Direct Lake mode:

1. **Line Chart:** plotted `avg_observed_cf`, `avg_historical_cf`, and `baseline_avg_cf` to compare current October performance with the 30-year baseline.
2. **Column Chart:** visualized `avg_deviation` by month, highlighting whether production was above or below expectations (−3.2 % for October 2025).
3. **Area Band Chart:** used `baseline_avg_cf ± baseline_std_cf` to illustrate natural historical variability, forming a shaded confidence band around the baseline trend.

Although the pilot currently covers ten days of Fingrid data (October 2025), the architecture is fully scalable: as new months are ingested, Power BI visuals will automatically extend. The project showcases the integration of real-time cloud ingestion, PySpark transformations, and business-ready analytics in Microsoft Fabric. It bridges energy-sector domain data and modern engineering practices—turning raw hydropower measurements into intuitive, evidence-based operational intelligence.

From **ETL → analytics → AI augmentation**:

Component	Purpose
Data Engineering workspace (Lakehouse)	Stores all gold-level Delta tables.
Semantic Model (Power BI / Fabric)	Defines relationships and measures (makes it queryable).
Lakehouse AI services or Copilot in Fabric Notebooks	For retrieval + LLM generation.
Vector Index (OneLake AI / Azure Cognitive Search)	Stores embeddings of time-series chunks or metadata for fast retrieval.

Since this is just 1 ECTS course, I did not go on refining the visuals in Power BI, however, I tried to demonstrate a fully working end-to-end analytical workflow following data fetching via a unified platform, data processing using medallion architecture and analytics in Power BI.

After this, I also wanted to try out generative AI capabilities, so I wanted to implement RAG in the time-series data that I had. I performed following things:

- **Build Feature Store:** Create a new Lakehouse table with time windows (hour/day) and derived features (CF mean, deviation, trend slope).
- **Generate Embeddings:** Could not find Fabric's *AI services* that can create *Text embeddings* on textualized summaries of each record. So proceeded with FAISS locally.
- **Create Vector Index:** Again, could not register the embeddings in Fabric's AI Catalog or Azure Cognitive Search. So created embeddings in the fabric notebook instead.
- **Create a Semantic Q&A Model:** Now, if I had succeeded in utilizing the Fabric's embeddings and vector index, I could have used Fabric Copilot or Azure OpenAI Service with retrieval plugin configured for the vector index. To substitute this, I wanted to implement this via code in the notebook.
- **Integrate into Dashboard:** Add a **chat visual** in Power BI (Copilot for Power BI) to query natural-language questions directly on top of our gold data. Now, this step was also unsuccessful.

Successful implementation could have had me an **intelligent hydropower assistant** inside Fabric that can:

- Answer questions like “*Show similar performance drops in past years.*”
- Explain anomalies in plain language, grounded in our dataset.
- Generate automated narrative summaries for monthly reports.

All built entirely within the Fabric workspace — using **existing gold tables** as the retrieval base, without moving data outside OneLake. Even though I did not succeed, I have documented my learnings.

Phase 1 — RAG Implementation Plan

Step 1 – Prepare the Base Data (Lakehouse)

Goal: ensure our Gold tables are clean, timestamp-indexed, and semantically rich.

Actions


- Use existing:
 - gold_cf_deviation (hour-level)
 - gold_cf_deviation_monthly (month-level)
 - gold_zenodo_baseline_monthly (baseline reference)
- Add a descriptive text column for each record, e.g. in a PySpark notebook.

After Phase 1, I had a **prototype RAG system** inside Fabric that:

- Accepts natural-language questions about hydropower performance.
- Retrieves matching time-series slices from our own Lakehouse.
- Generates context-aware explanations and anomaly reports.

All computation stays within OneLake → no external storage or pipelines needed.

Step 2 – Create an AI Workspace / AI Catalog in Fabric

Table	Purpose	Include in RAG?
gold_cf_deviation	 Main dataset — observed, historical, deviation (hour-level)	Yes → use this as df
gold_cf_deviation_monthly	Aggregated by month	Optional — for summary or trend explanations
gold_zenodo_baseline_monthly	Long-term climatology baseline	Optional — for background context if we want to retrieve 30-year norms

Created `summary_text` column that was purely numeric. By bringing metadata, I was able to compute national level context for hydropower. Thereafter, I was also able to add seasonal / baseline context using appropriate data source and build enriched summary text. Finally, the **retriever** was able to find relevant time periods *and* understand seasonal and typological context.

Step 3 – Generate Embeddings for Each Record

After this, I converted each record's summary_text into a numerical vector so it can be semantically searched.

Step 4 – Register a Vector Index (keycolumn timestamp)

Step 5 – Create an LLM Endpoint for Retrieval + Generation

Step 6 – Integrate into Power BI / Copilot Experience

These can now be iterated on various use cases:

Scenario	Implementation Path
Q&A Bot for Specific Times	Query = timestamp filter → retrieve hour/day records → summarize.
Pattern Search & Explanation	Retrieve top-K similar embeddings → describe similar historical episodes.
Anomaly Reporting	Detect z-score > 3 → send to LLM for narrative report generation.

While I tried Steps 4 – 6 to accomplish various scenarios/use-cases, I could not succeed. However, I have summarized my learnings as below:

Attempted RAG Implementation in Fabric

The extended objective was to **embed descriptive summaries** of each observation for *semantic retrieval* and *AI-driven insights*.

Implemented steps:

1. **Generated textual summaries** per observation (combining capacity, deviation, and month context).
2. **Created vector embeddings** using Azure OpenAI's text-embedding-3-small model.
3. **Indexed vectors** with FAISS for similarity search on historical-like events.

Encountered challenges:

- Fabric's **AI Foundry deployment region restrictions** (NorwayEast unavailable) prevented Chat model deployment.
- **Schema mismatches and serialization issues** arose when saving embedding arrays to Delta format.
- Attempted joins between FAISS results and metadata failed due to **index-type conflicts (long vs int)**.

These steps nonetheless demonstrated the **blueprint for an intelligent retrieval system** inside Fabric — where a user could query, e.g.,

“Show recent hours when Finland’s hydropower deviated most from baseline during autumn months,”

and the system would return semantically matched summaries from similar events.






5 What Would Have Followed

If the AI service were fully connected:

- A **chat interface** could summarize deviations and explain anomalies using retrieved context (“agentic” reasoning).
- The **Power BI dashboard** would include a natural-language Q&A visual for contextual search across `summary_text` and embeddings.

This would evolve the lakehouse from static analytics to a **self-explaining, data-aware assistant**.

6 Key Takeaways

-  Hands-on understanding of **Microsoft Fabric Lakehouse architecture**.
 -  Experience integrating **real-time + historical** energy datasets.
 -  Implementation of **Spark-based ETL and Delta management**.
 -  Exposure to embedding generation, vector indexing, and semantic retrieval workflows.
 -  Learned practical constraints of **region-restricted AI deployment** and schema consistency for embedding persistence.
-

In essence:

I successfully established the foundation for a *hydropower intelligence platform* — bridging data engineering, analytics, and AI.

While the RAG portion couldn’t be completed due to Fabric’s AI region limits, my design demonstrates a scalable path to **agentic time-series analytics**, where natural-language exploration meets data-driven hydropower insights.